

# Toy Models of Superposition

Tristan Thomas

November 2025

## Paper Introduction/Key Questions

In the paper<sup>1</sup> studied as part of this project, the authors build toy models of superposition.

Key questions:

- ▶ What is superposition?
- ▶ What is meant by toy model in this case?
- ▶ Is there extension from the "toy model" to describe effects in larger models?

---

<sup>1</sup>Nelson Elhage et al. *Toy Models of Superposition*. 2022. arXiv: 2209.10652 [cs.LG]. URL: <https://arxiv.org/abs/2209.10652>.


## What is superposition?

- ▶ Superposition is a phenomenon in which neural network models are able to represent more features than they have dimension.
- ▶ Take digit classification as an example.
  - ▶ Ideally, neurons are monosemantic
  - ▶ E.g. One neuron detects if there is a closed loop in the number, one detects straight edges, etc.
- ▶ Is this really what happens? Typically, no.

# The fuel for superposition-feature sparsity

As Sparsity Increases, Models Use "Superposition" To Represent More Features Than Dimensions

Increasing Feature Sparsity →




0% Sparsity

The two most important features are given **dedicated orthogonal dimensions**, while other features are **not embedded**.



80% Sparsity

The four most important features are represented as **antipodal pairs**. The least important features are **not embedded**.



90% Sparsity

All five features are embedded as a **pentagon**, but there is now "positive interference."

Feature Importance

- Yellow circle: Most important
- Green circle: Medium important
- Dark green circle: Least important

Figure: Example with training embedding 5 features in 2D space.

## Related Work

Mikolov et. al<sup>2</sup> Investigated the semantic properties of word representations. i.e. King + Male - Queen = Female.

There is a lot of work on polysemantic neurons. These neurons may respond to very unrelated words or concepts in language models (Clemson football and South African safari tours).<sup>3</sup>


---

<sup>2</sup>Tomas Mikolov et al. *Efficient Estimation of Word Representations in Vector Space*. 2013. arXiv: 1301.3781 [cs.CL]. URL: <https://arxiv.org/abs/1301.3781>.


<sup>3</sup>Chris Olah, Alexander Mordvintsev, and Ludwig Schubert. “Feature Visualization”. In: *Distill* (2017). <https://distill.pub/2017/feature-visualization>. DOI: 10.23915/distill.00007.

# The Superposition Hypothesis

- ▶ Non-privileged basis: Features are embedded in any direction
- ▶ Privileged basis: Features are incentivized to align with basis dimensions
- ▶ Superposition says the network wants to represent more features than there are neurons. This is mathematically justified by the Johnson-Lindenstrauss Theorem.




**Polysemanticity** is what we'd expect to observe if features were not aligned with a neuron, despite incentives to align with the privileged basis.




In the **superposition hypothesis**, features can't align with the basis because the model embeds more features than there are neurons. Polysemanticity is inevitable if this happens.

# Issues



Even if only **one sparse feature** is active, using linear dot product projection on the superposition leads to **interference** which the model must tolerate or filter.



If the features aren't as sparse as a superposition is expecting, **multiple present features** can additively interfere such that there are multiple possible nonlinear reconstructions of an **activation vector**.

- ▶ Setup: To understand the left figure, we define a concrete example with 3 features (Universities) packed into 2 neurons.
- ▶ Feature Directions: We assign a vector in 2D space for each university:
  - ▶ Clemson:  $v_1 = \left(\frac{\sqrt{3}}{2}, \frac{1}{2}\right)$
  - ▶ Duke:  $v_2 = \left(-\frac{\sqrt{3}}{2}, \frac{1}{2}\right)$
  - ▶ NC State:  $v_3 = (0, -1)$
- ▶ Weight Matrix ( $W$ ): Stacking these as columns gives us our encoding matrix:

$$W = \begin{pmatrix} \frac{\sqrt{3}}{2} & -\frac{\sqrt{3}}{2} & 0 \\ \frac{1}{2} & \frac{1}{2} & -1 \end{pmatrix}$$

## Quantifying Interference

- ▶ In low dimensional space, we can't have all three vectors be even close to mutually orthogonal.
- ▶ We project the Duke vector onto the Clemson vector:

$$\begin{aligned}\text{proj}_C(D) &= \underbrace{\left( \begin{pmatrix} -\frac{\sqrt{3}}{2} \\ \frac{1}{2} \end{pmatrix} \cdot \begin{pmatrix} \frac{\sqrt{3}}{2} \\ \frac{1}{2} \end{pmatrix} \right)}_{\text{Dot Product}} \underbrace{\begin{pmatrix} \frac{\sqrt{3}}{2} \\ \frac{1}{2} \end{pmatrix}}_{\text{Direction}} \\ &= \left( -\frac{3}{4} + \frac{1}{4} \right) \text{Clemson} = -0.5 \text{Clemson}\end{aligned}$$

- ▶ Result: It is obvious from this derivation that activating Duke will have an effect on the apparent activation of Clemson.

## Ambiguity

- ▶ Activate Clemson and Duke both at **0.9**.
- ▶ 
$$\begin{pmatrix} \frac{\sqrt{3}}{2} & -\frac{\sqrt{3}}{2} & 0 \\ \frac{1}{2} & \frac{1}{2} & -1 \end{pmatrix} \begin{pmatrix} 0.9 \\ 0.9 \\ 0 \end{pmatrix} = \begin{pmatrix} 0 \\ 0.9 \end{pmatrix}$$
- ▶ From previous results, if you de-embed this vector you will get half the Clemson feature value (because the Duke projects on and cancels half).
- ▶ NC State Result: How does this look to the NC State feature  $\begin{pmatrix} 0 \\ -1 \end{pmatrix}$ ?

$$\text{Activation} = \begin{pmatrix} 0 \\ 0.9 \end{pmatrix} \cdot \begin{pmatrix} 0 \\ -1 \end{pmatrix} = -0.9$$

- ▶ Conclusion: The presence of Clemson and Duke creates an ambiguous signal of -0.9 on the NC State channel.

# The Full Dembedding Math

- ▶ Recovery ( $x' = W^T h$ ): We get the activation of each feature by multiplying our activation vector by the transpose of the feature matrix.

$$\begin{bmatrix} x'_{\text{Clemson}} \\ x'_{\text{Duke}} \\ x'_{\text{NC State}} \end{bmatrix} = \underbrace{\begin{bmatrix} \frac{\sqrt{3}}{2} & \frac{1}{2} \\ -\frac{\sqrt{3}}{2} & \frac{1}{2} \\ 0 & -1 \end{bmatrix}}_{W^T \text{ (Feature Rows)}} \underbrace{\begin{bmatrix} 0 \\ 0.9 \end{bmatrix}}_{\text{Activation } h} = \begin{bmatrix} 0.45 \\ 0.45 \\ -0.9 \end{bmatrix}$$

- ▶ Results:
  - ▶ Clemson & Duke: Recovered at 0.45
  - ▶ NC State: Introduced a nonzero value (-0.9).
- ▶ Interpretation: We now have no idea what features were really present at the input. A problem due to not having sparsity.

# Testing the Superposition Hypothesis

- ▶ Can we project a vector from high dimensional space to low dimensional space and then recover it?
- ▶ Data is based on each  $x_i$  having a sparsity  $S_i$  and important  $I_i$ .  $x_i = 0$  with probability  $S_i$ , otherwise uniformly sampled from  $[0, 1]$ .

HYPOTHETICAL DISENTANGLING MODEL



Our first experiments will test the extent to which the idealized activations of an imagined larger model can be **stored** and **recovered** from a lower-dimensional space.

OBSERVED MODEL



## Two Models to Consider

- ▶ Linear model
  - ▶  $h = Wx$
  - ▶  $x' = W^T h + b$
  - ▶  $x' = W^T Wx + b$
- ▶ ReLU Output Model
  - ▶  $h = Wx$
  - ▶  $x' = \text{ReLU}(W^T h + b)$
  - ▶  $x' = \text{ReLU}(W^T Wx + b)$
- ▶ Note that we add new things to the models from the previous illustrative example. We have a bias that can help put features at their expected value (if we know they should be sparse). We also add activation.

# Results Visualization Example



It tends to be easier to visualize  $W^T W$  than  $W$ .  
Here we see that  $W^T W$  is an **identity matrix** for the most important features and **0** for less important ones.



We can also look at the bias,  $b$ .  
The bias is **zero** for features learned to pass through, and the **expected value** (a positive number) for others.



**Figure:** If a feature is represented, it should have a norm close to 1. The biases will be zero if we have a representation of the feature, otherwise it can be used to help us get a closer representation for features not well represented by embedding.

# Results

- ▶ Will show them directly in paper for better viewing experience
- ▶ Conclusion:
  - ▶ The linear model only give priority to the 5 most important features
  - ▶ As we look at ReLU models, they are still only embedding important features as sparsity being so high will damage feature recovery
  - ▶ As sparsity is increased, more features are stored as interference and ambiguous effects are mitigated

# Citations

Figures were borrowed from [1]